



# AutoML for Efficient Vision Learning



Mingxing Tan

Google Brain, joint work with many Google colleagues


02/26/2021


Including results by many other people at Google and elsewhere.

# Deep learning is taking off ...




**But, designing DL models is hard...**








LeCun and Bengio 1995.  
Convolutional networks for  
images, speech, and time-series.




Krizhevsky et al 2012. Imagenet classification with deep  
convolutional neural networks.



Huang et al 2017.  
Densely connected  
convolutional networks



Szegedy et al 2016. Rethinking the inception architecture for computer  
vision




Szegedy et al 2017  
Inception-v4,  
Inception-Resnet, and the  
Impact of Residual  
Connections on Learning.




+others


Szegedy et al. 2015.  
Going deeper with  
convolutions.



He et al. 2016. Deep  
residual learning for  
image recognition.



Xie et al 2017. Aggregated  
residual transformations for  
deep neural networks



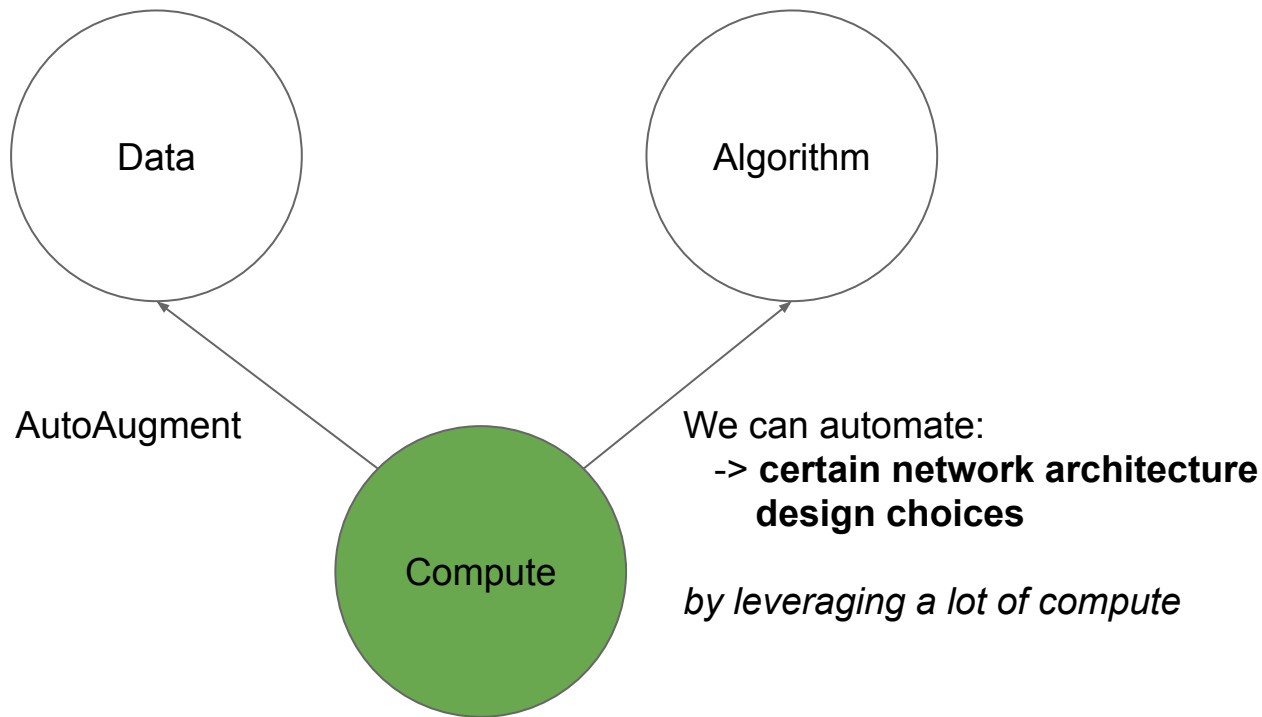
Hu et al. 2018. Squeeze-and-Excitation Networks

Can we automate deep learning? AutoML?

**AUTOMATE**



# NAS: Neural Architecture Search



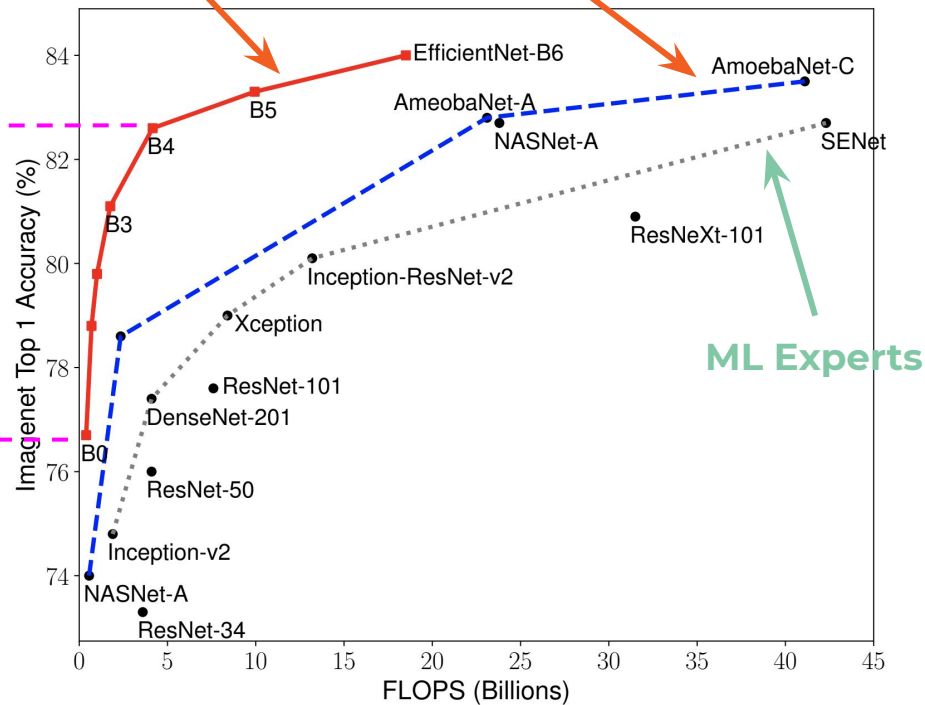
# Image Recognition

AutoML 2019

AutoML 2017

Step2: Scale it up

Step1: find a good model with NAS



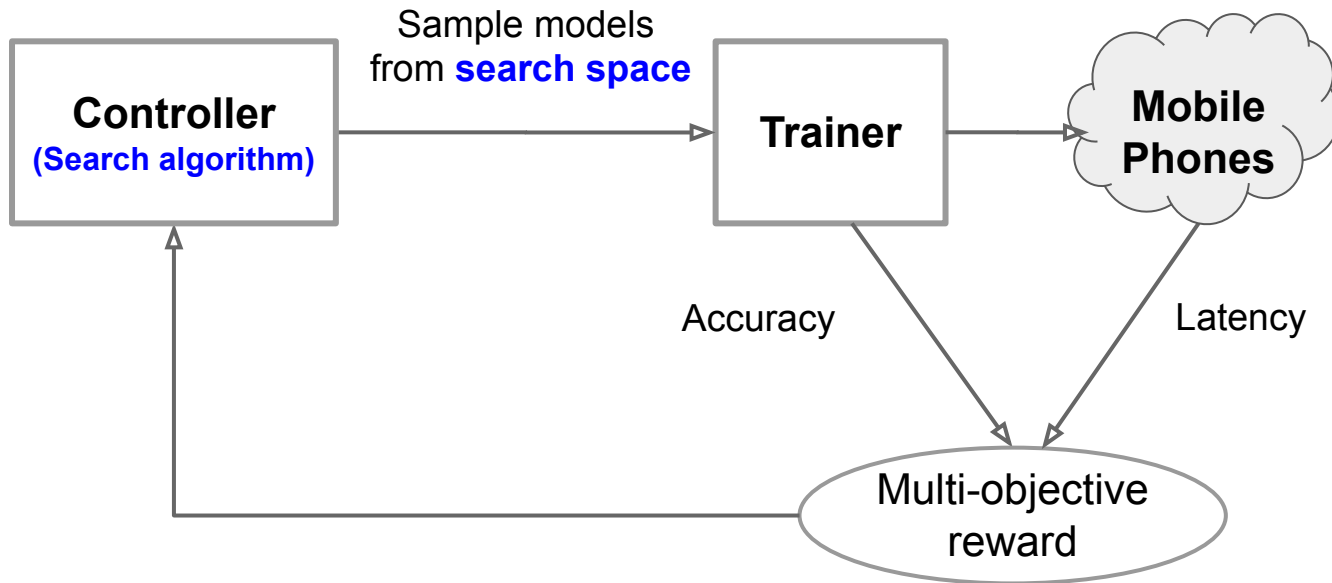
# MNAS: Platform-Aware Neural Architecture Search

Step1: find a good compact model with NAS

Mingxing Tan, et. al. *MnasNet: Platform-Aware Neural Architecture Search for Mobile*. CVPR 2019. <https://arxiv.org/abs/1807.11626>


Mingxing Tan, Quoc V. Le *MixNet: Mixed Depthwise Convolutional Kernels*. BMVC 2019. <https://arxiv.org/abs/1907.09595>

Andrew Howard, et al. *MobileNetV3: Searching for MobileNetV3*. ICCV 2019. <https://arxiv.org/abs/1905.02244>

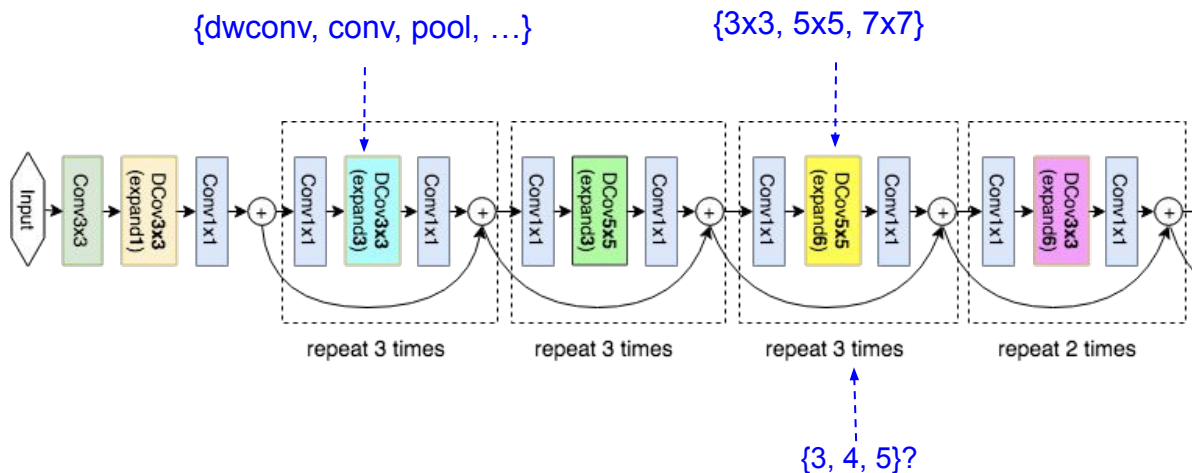




# Search Space: A cell or a network?




Search for a Cell  
And repeatedly stack the cell



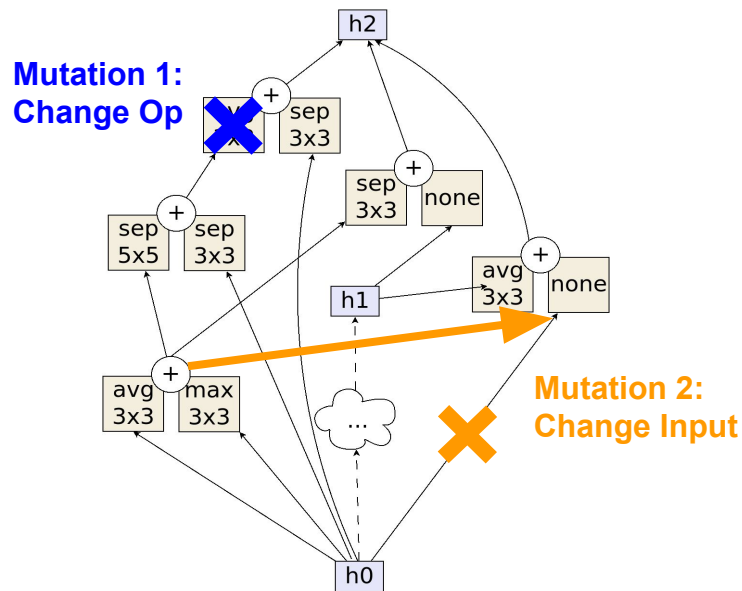
Search for a network

# Search Algorithm: RL or Evolution



NAS with Reinforcement Learning [ICLR'17]

Flexible & Expensive




NAS with Evolution [AAAI'19]

Flexible & Expensive

10,000 GPU hours?



# The Rescue: Weight Sharing



## Key idea:

1. One path inside a big model is a child model
2. Controller selects a path inside a big model and train for a few steps
3. Controller selects another path inside a big model and train for a few steps, reusing the weights produced by the previous step
4. Etc.

Results: **Can save 100->1000x compute**

Related work: DARTS, SMASH, One-shot, ProxylessNAS, FBNet, ...

# Multi-Object Reward

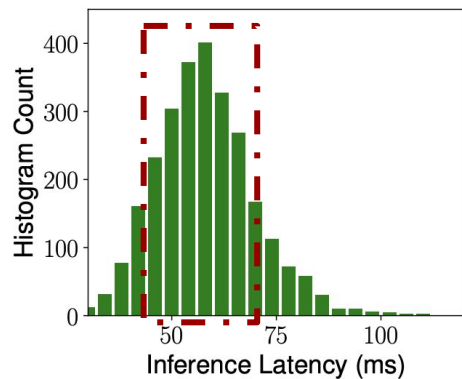

- Incorporate both **latency and accuracy** into reward function
  - *Weighted product* to approximate Pareto-Optimal solutions

$$\underset{m}{\text{maximize}} \quad ACC(m) \times \left[ \frac{LAT(m)}{T} \right]^w \quad (2)$$

where  $w$  is the weight factor defined as:



$$w = \begin{cases} \alpha, & \text{if } LAT(m) \leq T \\ \beta, & \text{otherwise} \end{cases} \quad (3)$$

Latency as **hard** constraint




(a)  $\alpha = 0, \beta = -1$

Latency as **soft** constraint



(b)  $\alpha = \beta = -0.07$

# Example Model





**Observation 1:**

5x5 depconv is very common

**Observation 2:**

Diverse layers throughout the network

# ImageNet Results

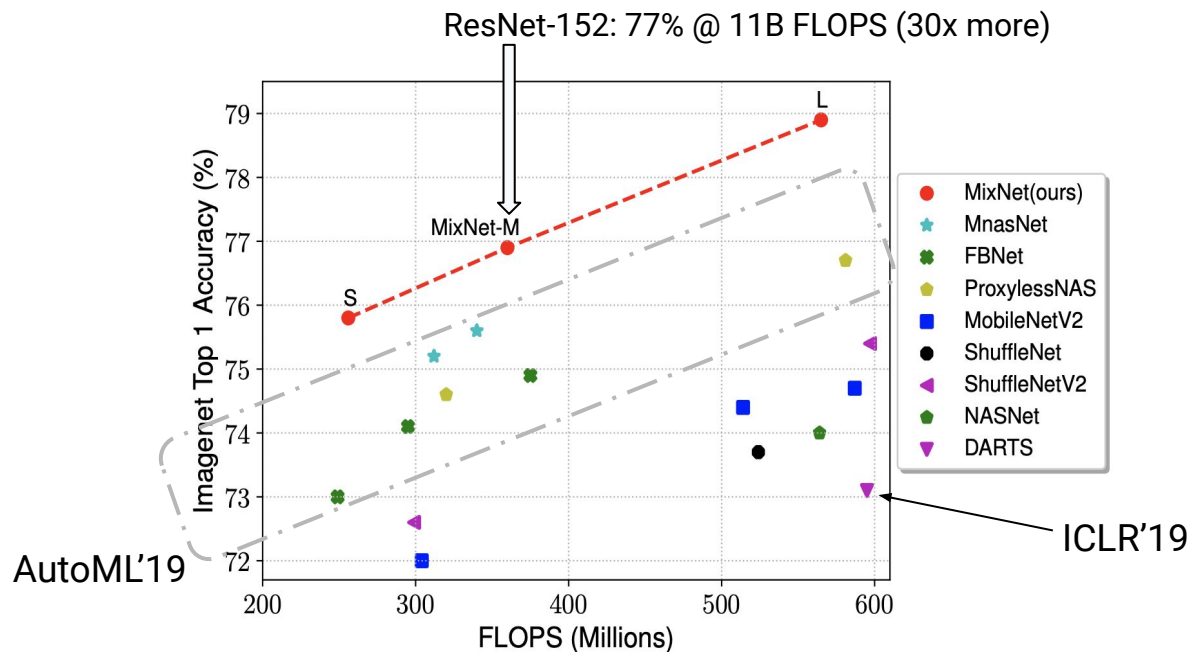



Mingxing Tan, et. al. *MnasNet: Platform-Aware Neural Architecture Search for Mobile*. CVPR 2019. <https://arxiv.org/abs/1807.11626>

Andrew Howard, et al. *MobileNetV3: Searching for MobileNetV3*. ICCV 2019. <https://arxiv.org/abs/1905.02244>

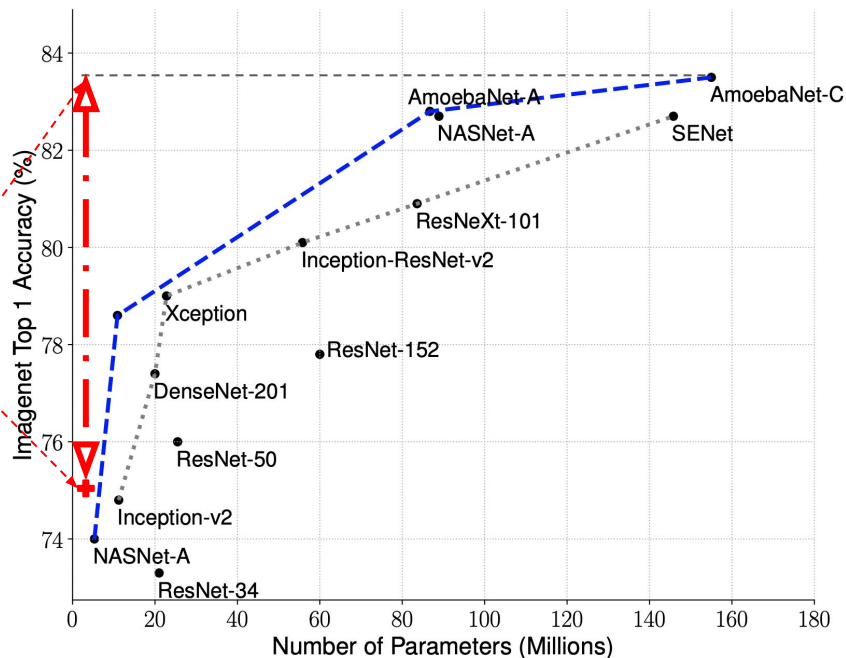


# Use NAS as a tool: MNAS -> MixMet



# So far so good, but wait ...

Huge accuracy gap to SoTA



**NOT SATISFIED**



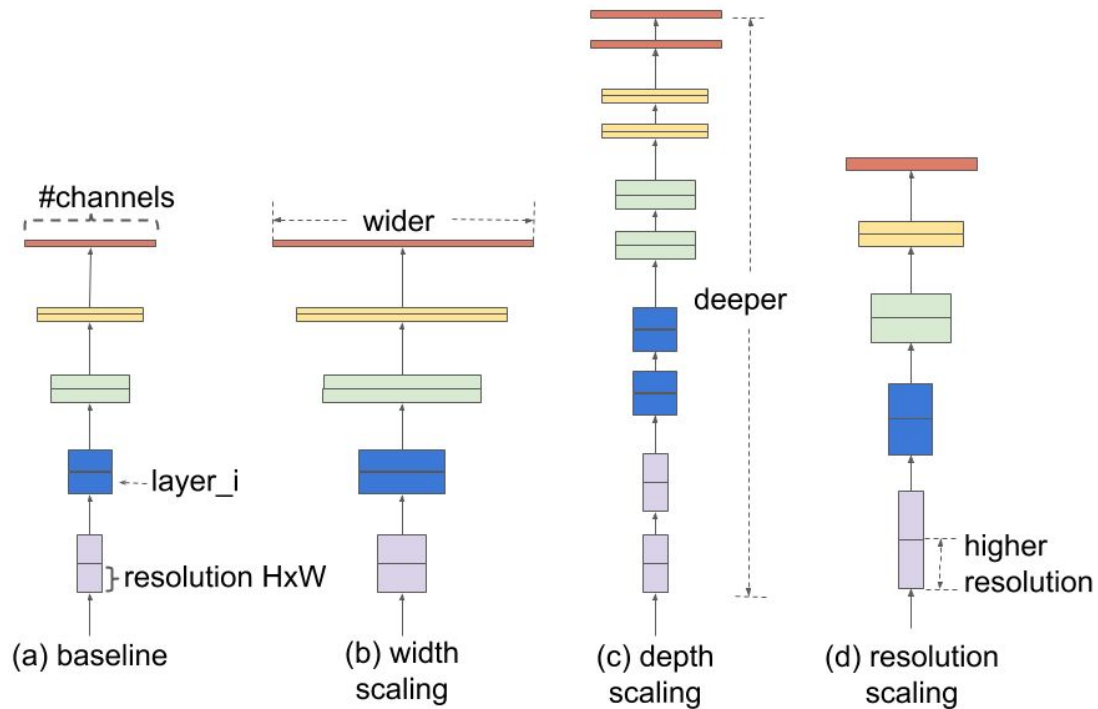
# EfficientNet/Det: Scaling Up ConvNets

Step2: towards better accuracy & efficiency




Mingxing Tan, Quoc V. Le. *EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks*. ICML 2019. <https://arxiv.org/abs/1905.11946>

Mingxing Tan, Ruoming Pang, Quoc V. Le. *EfficientDet: Scalable and Efficient Object Detection*. CVPR 2020. <https://arxiv.org/abs/1911.09070>

# How to Scale Up A ConvNet?




# Limitations of Single-Dimension Scaling



Accuracy saturates quickly if scaling by any single dimension

# How to Scale Up A ConvNet?



# Compound Scaling

depth:  $d = \alpha^\phi$

width:  $w = \beta^\phi$

resolution:  $r = \gamma^\phi$

s.t.  $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$

$\alpha \geq 1, \beta \geq 1, \gamma \geq 1$

## Step1:

- First fix  $\phi = 2$ , and find  $\alpha, \beta, \gamma$  with local search.


## Step2:

- Then fix  $\alpha, \beta, \gamma$ , and scale the network with different  $\phi$ .


Compound scaling improves MobileNetV1, MobileNetV2, and ResNet-50.



# Scaling the Same Baseline EfficientNet-B0


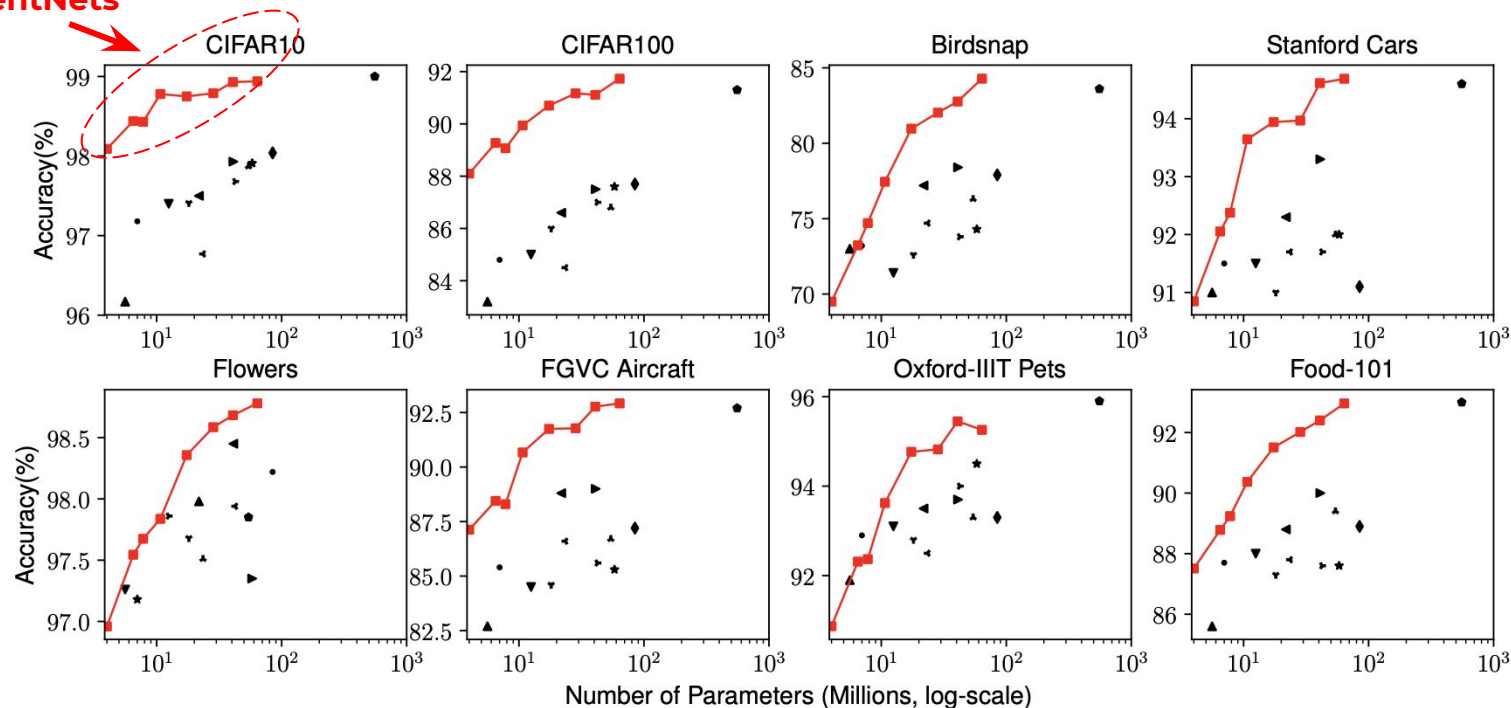


# ImageNet Results




# Transfer Learning Results


**EfficientNets**



# Object Detection




**EfficientDet: a new family of detection models based on EfficientNet backbones**



# Others

## Semantic Segmentation (Cityscapes)



## Point Cloud 3D Detection (Waymo OD):

Metric: APH/L2	PED	VEH	CYC	ALL_NS
ShyPillars Single	<b>0.6884</b>	<b>0.6623</b>	<b>0.6723</b>	<b>0.6743</b>
Prev. Best Single	0.6170	0.6446	0.6204	0.6273


## Video Action Recognition (Kinetics600):

**Accuracy > X3D, but**

- 67% fewer FLOPs
- 65% less memory.

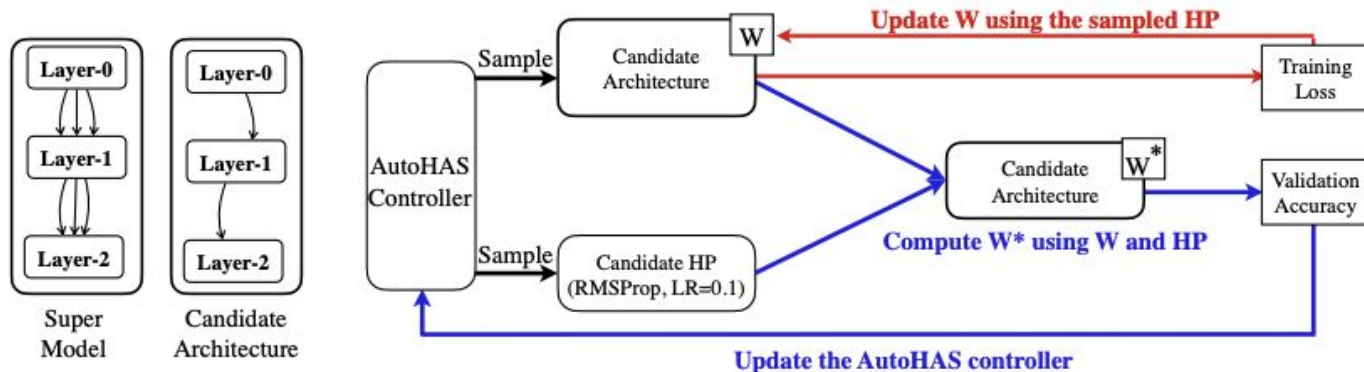
# **Wild NAS Topics**

How about hyperparamters?






# AutoHAS: Hyperparameter and Architecture Search



*Figure 1. The overview of AutoHAS.* LEFT: Each candidate architecture's weights are shared with a super model, where each candidate is a sub model within this super model. RIGHT: During the search, AutoHAS alternates between optimizing the shared weights of super model  $\mathcal{W}$  and updating the controller. It also creates temporary weights  $\mathcal{W}^*$  by optimizing the sampled candidate architecture using the sampled candidate hyperparameter (HP). This  $\mathcal{W}^*$  will be used to compute the validation accuracy as a reward so as to update the AutoHAS controller to select better candidates. Finally,  $\mathcal{W}^*$  is discarded after updating the controller so as not to affect the original  $\mathcal{W}$ .

**Key idea: Jointly search for neural network and hyperparameters**

# AutoHAS: Hyperparameter and Architecture Search



**10x faster than Vizier**

ImageNet	77.2% -> 77.9% (+0.7%)
CIFAR-100	76.3% -> 78.4% (+2.1%)
Flowers	74.0% -> 85.4% (+11.4%)

**Accuracy gains on different datasets**

Can we use NAS to search for back-propagation?



# AutoML Zero: search for models with basic math ops

## (1) Programs with 3 functions

```
def Setup():  
def Predict():  
def Learn():
```

Scalars: s0, s1, ...  
Vectors: v0, v1, ...  
Matrices: m0, m1, ...

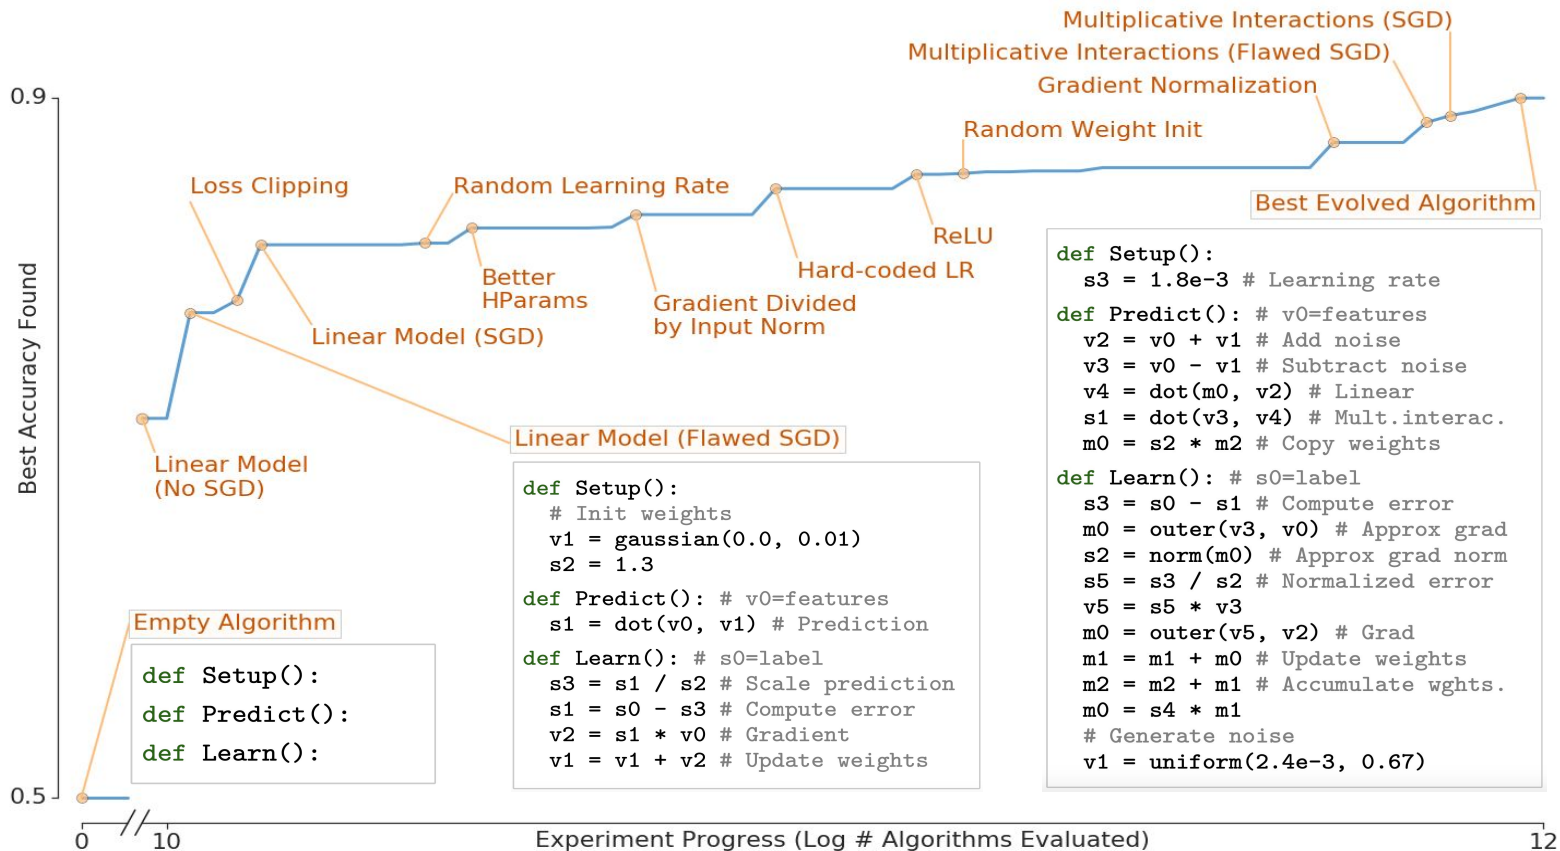
## (2) Basic math operations (64)

- Arithmetics
  - $a + b$ ,  $a * b$ , ...
- Trigonometry
  - $\sin(a)$ ,  $\cos(a)$ , ...
- Pre-calculus
  - $\exp(a)$ ,  $\log(a)$ , ...
- Linear algebra
  - $\text{dot}(a, b)$ ,  $\text{norm}(a)$ , ...
- Probability & Stats
  - $\text{gaussian}(a, b)$ ,  $\text{mean}(a)$ , ...

## (3) Evaluation

```
# (Setup, Predict, Learn) is the input ML algorithm.  
# Dtrain / Dvalid is the training / validation set.  
# sX/vX/mX: scalar/vector/matrix var at address X.  
def Evaluate(Setup, Predict, Learn, Dtrain, Dvalid):  
    # Zero-initialize all the variables (sX/vX/mX).  
    initialize_memory()  
    Setup() # Execute setup instructions.  
  
    for (x, y) in Dtrain:  
        v0 = x # x will now be accessible to Predict.  
        Predict() # Execute prediction instructions.  
        # s1 will now be used as the prediction.  
        s1 = Normalize(s1) # Normalize the prediction.  
        s0 = y # y will now be accessible to Learn.  
        Learn() # Execute learning instructions.  
  
    sum_loss = 0.0  
    for (x, y) in Dvalid:  
        v0 = x  
        Predict() # Only execute Predict(), not Learn().  
        s1 = Normalize(s1)  
        sum_loss += Loss(y, s1)  
  
    mean_loss = sum_loss / len(Dvalid)  
    # Use validation loss to evaluate the algorithm.  
    return mean_loss
```

# AutoML Zero



**Thank you**